

Classification Level: Top Secret() Secret() Internal() Public(✓)

RKLLM SDK User Guide

(Graphic Computing Platform Center)

Mark:	Version:	1.1.4
<input type="checkbox"/> Changing	Author:	AI Group
<input checked="" type="checkbox"/> Released	Completed Date:	7/February/2025
	Reviewer:	Vincent
	Reviewed Date:	7/February/2025

瑞芯微电子股份有限公司

Rockchip Electronics Co., Ltd.

(Copyright Reserved)

Revision History

Version	Modifier	Date	Modify description	Reviewer
V1.0.0	AI Group	2024-3-23	Initial version	Vincent
V1.0.1	AI Group	2024-5-8	<ol style="list-style-type: none"> 1. Optimize memory usage. 2. Optimize inference time. 3. Optimize quantization accuracy. 4. Support Gemma, Phi-3 and other models. 5. Add Server call and interrupt interface. 	Vincent
V1.0.1	AI Group	2024-10-10	<ol style="list-style-type: none"> 1. Support group-wise quantization. 2. Support joint inference with lora model loading. 3. Support storage and preloading of prompt cache. 4. Optimize initialization, prefill, and decode time. 5. Support gguf model conversion. 6. Add gdq algorithm to improve 4-bit quantization accuracy. 7. Add mixed quantization algorithm, supporting a combination of grouped and non-grouped quantization based on specified ratios. 8. Add support for models such as Llama3, Gemma2, and Minicpm3. 	Vincent
V1.1.4	AI Group	2025-02-07	<ol style="list-style-type: none"> 1. Support for GPTQ model conversion 2. Added support for TeleChat2 and MiniCPM-S models 3. Added update_rkllm interface 4. Support for LLM model conversion in MiniCPMV and Qwen2VL 5. Fixed inference issues in MiniCPM3 6. Added DeepSeek_R1_Distill_Demo 7. Added Qwen2_VL_2B_Demo 	Vincent

Table of contents

1 Introduction to RKLLM.....	5
1.1 Introduction to RKLLM Toolchain.....	5
1.1.1 Introduction to RKLLM-Toolkit.....	5
1.1.2 Introduction to RKLLM Runtime	5
1.2 Introduction to the RKLLM Development Process	5
1.3 Applicable Hardware Platforms.....	7
2 Development Environment Preparation	8
2.1 Installation of RKLLM-Toolkit	8
2.1.1 Installation via Pip.....	9
2.2 Introduction for RKLLM Runtime Usage.....	10
2.3 Compilation Requirements for RKLLM Runtime	10
2.4 Kernel Update.....	10
3 RKLLM Instruction Guide	13
3.1 Model Conversion.....	13
3.1.1 RKLLM Initialization.....	13
3.1.2 Loading Models	13
3.1.3 Construction for RKLLM Model.....	14
3.1.4 Exporting Model	16
3.1.5 GPTQ Model Conversion.....	17
3.1.6 Update For Old Version Model	18
3.1.7 Simulation Accuracy Evaluation.....	19
3.1.8 Simulated Model Inference.....	20
3.2 Inference Implementation in Board-side.....	21
3.2.1 Define Callback Function	21
3.2.2 Define RKLLMParam	23

3.2.3 Definite Input Struct.....	26
3.2.4 Initialize Model.....	30
3.2.5 Inference Model.....	31
3.2.6 Interrupt Model Inference.....	32
3.2.7 Release Model.....	32
3.2.8 Load LoRA Model	33
3.2.9 Load Prompt Cache	34
3.3 Board-side Inference Example	36
3.3.1 Complete Code of Example Project.....	36
3.3.2 Instructions for Example Project.....	36
3.3.3 Monitor inference performance	37
3.4 Implementation of Board-side Server	37
3.4.1 Deployment Example of RKLLM-Server-Flask.....	38
3.4.2 Deployment Example of RKLLM-Server-Gradio.....	45
4 Reference.....	51

1 Introduction to RKLLM

1.1 Introduction to RKLLM Toolchain

1.1.1 Introduction to RKLLM-Toolkit

RKLLM-Toolkit is a development kit that provides users with the ability to quantify and convert large language models on the PC. Through the Python interface provided by the tool, the following functions can be conveniently accomplished:

1) Model Conversion: It supports the conversion of Large Language Model (LLM) in Hugging Face or GGUF format to RKLLM model, and the supported models include LLaMA, Qwen, Qwen2, Phi-2, Phi-3, ChatGLM3, Gemma, Gemma2, InternLM2, TeleChat2, MiniCPM-S, MiniCPM and MiniCPM3. The converted RKLLM models can be loaded and used on Rockchip NPU platform.

2) Quantization Function: Supports quantization of floating-point models into fixed-point models. Currently supported quantization types include:

- a. w4a16;
- b. w4a16 grouped quantization (supported group sizes: 32, 64, 128);
- c. w8a8;
- d. w8a8 grouped quantization (supported group sizes: 128, 256, 512);

1.1.2 Introduction to RKLLM Runtime

The RKLLM Runtime is primarily responsible for loading the RKLLM models obtained from the RKLLM-Toolkit conversion and executing inference via the Rockchip NPU drivers embedded in the RK3576/RK3588 series. Throughout the RKLLM model inference process, users retain the autonomy to specify inference parameter configurations, define various text generation methodologies, and seamlessly retrieve inference results from the model via pre-defined callback functions.

1.2 Introduction to the RKLLM Development Process

The overall development process of RKLLM primarily comprises two phases: model conversion in

PC and deployment and execution in board-side:

1) Model Conversion in PC:

During this phase, the provided Hugging Face formatted LLM is converted into RKLLM format for efficient inference on the Rockchip NPU platform. This step includes:

- a. Model Acquisition: Obtain the large language model from:
 - a) Open-source large language models in Hugging Face format.
 - b) Self-trained large language models, with the requirement that the saved model structure is consistent with models on the Hugging Face platform.
 - c) GGUF models, currently supporting only q4_0 and fp16 types.
- b. Model Loading: Load the model in huggingface format with the `rkllm.load_huggingface()` function and load the gguf format model with `rkllm.load_gguf()` function;
- c. Model Quantization Configuration: Utilize the `rkllm.build()` function to construct the RKLLM model, with the option to perform model quantization for enhanced hardware deployment performance. Additionally, choose different optimization levels and quantization types during the construction process.
- d. Model Export: Export the RKLLM model as a `.rkllm` format file using the `rkllm.export_rkllm()` function for subsequent deployment.

2) Deployment and Execution in Board-side:

This phase encompasses the actual deployment and execution of the model, typically involving the following steps:

- a. Model Initialization: Load the RKLLM model onto the Rockchip NPU platform, configure model parameters accordingly to define the desired text generation methods, and pre-define callback functions to receive real-time inference results, preparing for inference.
- b. Model Inference: Execute the inference operation by passing input data to the model and running model inference. Users can continuously retrieve inference results through pre-defined callback functions.

c. Model Release: After completing the inference process, release the model resources to allow other tasks to utilize the computational resources of the NPU.

These two steps constitute the complete RKLLM development process, ensuring successful conversion, debugging, and ultimately efficient deployment of the LLMs on the Rockchip NPU.

1.3 Applicable Hardware Platforms

The hardware platforms to which this document applies mainly include: RK3576, RK3588

2 Development Environment Preparation

The RKLLM toolchain zip file contains the whl installer for the RKLLM-Toolkit, the associated files of the RKLLM Runtime library and reference sample code. The specific folder structure is shown below:

```
doc
├── Rockchip_RKLLM_SDK_CN.pdf      # RKLLM SDK Documentation (Chinese)
└── Rockchip_RKLLM_SDK_CN.pdf      # RKLLM SDK Documentation (English)

examples
├── DeepSeek-R1-Distill-Qwen-1.5B_Demo # Board-side Inference Example
├── Qwen2-VL-2B_Demo      # Multimodal Inference Example
└── rkllm_server_demo      # RKLLM-Server Deployment Example

rkllm-runtime
├── runtime
│   ├── Android
│   │   └── librkllm_api
│   │       └── arm64-v8a
│   │           ├── librkllmrt.so # RKLLM Runtime Library
│   │           └── include
│   │               └── rkllm.h      # Header File
│   └── Linux
│       ├── librkllm_api
│       │   └── aarch64
│       │       ├── librkllmrt.so # RKLLM Runtime Library
│       │       └── include
│       │           └── rkllm.h      # Header File
└── rkllm-toolkit
    ├── rkllm_toolkit-x.x.x-cp38-cp38-linux_x86_64.whl
    └── rkllm_toolkit-x.x.x-cp310-cp310-linux_x86_64.whl

rknpu-driver
└── rknpu_driver_x.x.x_XXXXXXX.tar.bz2

scripts
├── fix_freq_rk3576.sh # Fixed-frequency Script for RK3576
└── fix_freq_rk3588.sh # Fixed-frequency Script for RK3588
```

This chapter provides detailed instructions for installing the RKLLM-Toolkit and RKLLM Runtime.

For specific usage instructions, please refer to the instructions in Chapter 3.

2.1 Installation of RKLLM-Toolkit

This section mainly describes how to install the RKLLM-Toolkit with pip install command. Users can refer to the following detailed instructions to complete the installation of the RKLLM-Toolkit toolchain.

2.1.1 Installation via Pip

2.1.1.1 Installation of Miniforge3

To avoid the need for multiple versions of Python environments, it is recommended to use miniforge3 to manage your Python environments.

Check if miniforge3 is installed and the version information of conda. If it is already installed, you can skip this section.

```
conda -V
# If "conda: command not found" is displayed, it indicates that conda
is not installed.
# For example, version information could be displayed as conda 23.9.0
```

Download the miniforge3 installation package.

```
wget -c https://mirrors.bfsu.edu.cn/github-release/conda-
forge/miniforge/LatestRelease/Miniforge3-Linux-x86_64.sh
```

Install miniforge3.

```
chmod 777 Miniforge3-Linux-x86_64.sh
bash Miniforge3-Linux-x86_64.sh
```

2.1.1.2 Create RKLLM-Toolkit Conda Environment

Activate the Conda base environment.

```
source ~/miniforge3/bin/activate # directory of miniforge3
# (base) xxx@xxx-pc:~$
```

Create a Conda environment named RKLLM-Toolkit with Python version 3.8 (recommended).

```
conda create -n RKLLM-Toolkit python=3.8
```

Activate the RKLLM-Toolkit Conda environment.

```
conda activate RKLLM-Toolkit
# (RKLLM-Toolkit) xxx@xxx-pc:~$
```

2.1.1.3 Installing RKLLM-Toolkit

In the RKLLM-Toolkit Conda environment, use the pip command to install the provided toolchain .whl package directly. During the installation process, the installer will automatically download the necessary dependencies for the RKLLM-Toolkit tools.

```
pip3 install rkllm_toolkit-x.x.x-cp38-cp38-linux_x86_64.whl
Or
pip3 install rkllm_toolkit-x.x.x-cp310-cp310-linux_x86_64.whl
```

If executing the following command does not result in any errors, the installation is successful.

```
python
from rkllm.api import RKLLM
```

2.2 Introduction for RKLLM Runtime Usage

Within the RKLLM toolchain files provided, the following components are included for the RKLLM runtime:

- 1) lib/librkllmrt.so: RKLLM Runtime library suitable for RK3576/RK3588 series.
- 2) include/rkllm.h: Corresponding header file to librkllmrt.so, containing descriptions of related structures and function definitions.

When constructing deployment inference code for RK3576/RK3588 series using the RKLLM toolchain, it is critical to ensure proper linkage to the above header file and function library to ensure compilation correctness. During the actual runtime of the code on RK3576/RK3588 boards, it is equally important to ensure successful transfer of the above function library files to the board and complete library declaration through the following environment variable settings:

```
export LD_LIBRARY_PATH=/path/to/your/lib
```

2.3 Compilation Requirements for RKLLM Runtime

When utilizing RKLLM Runtime, it's essential to pay attention to the version of the gcc compilation tool. We recommend the cross-compilation tool [gcc-arm-10.2-2020.11-x86_64-aarch64-none-linux-gnu](https://gcc.gnu.org/nightly/20201111-x86_64-aarch64-none-linux-gnu/). Please note that cross-compilation tools often have backward compatibility but not upward compatibility, so refrain from using versions below 10.2.

If you choose to use the Android platform, and you need to compile Android executable files, we recommend using the Android NDK tool for cross-compilation. You can download it from the following link: [Android NDK Cross-Compilation Tool Download Link](#). We recommend using version r21e.

You can also refer to the specific compilation methods in the compilation scripts located in the examples/DeepSeek-R1-Distill-Qwen-1.5B_Demo/deploy directory.

2.4 Kernel Update

Due to the requirement of a higher version of the NPU kernel for the provided RKLLM, users need

to verify that the NPU kernel on the board is version v0.9.8 before using RKLLM Runtime for model inference. The specific query command is as follows:

```
# Execute the following command to query the NPU kernel version.
cat /sys/kernel/debug/rknpu/version

# Confirm that the command output is as follows:
# RKNPU driver: v0.9.8
```

If the queried NPU kernel version is lower than v0.9.8, please proceed to the official firmware address to download the latest firmware for updating.

For users using non-official firmware, it's necessary to update the kernel. The RKNPU driver package supports two main kernel versions: kernel-5.10 and kernel-6.1. For kernel-5.10, it is recommended to use a specific version at [GitHub-rockchip-linux/kernelatdevelop-5.10](https://github.com/rockchip-linux/kernelatdevelop-5.10), and 5.10.209 is recommended. For kernel 6.1, it is recommended to use a specific version such as 6.1.84. Users can confirm the specific version number in the Makefile in the kernel's root directory. The specific steps to upgrade the kernel are as follows:

- 1) Download the [rknpu_driver_0.9.8_20241009.tar.bz2](#).
- 2) Unzip the compressed file and overwrite the RKNPU driver into the current kernel code directory.
- 3) Recompile the kernel.
- 4) Flash the newly compiled kernel to the device.

If an error log, as shown in Figure 2-1, is encountered during the kernel compilation process:

```
drivers/rknpu/rknpu_gem.c: In function 'rknpu_gem_mmap_pages':
drivers/rknpu/rknpu_gem.c:891:2: error: implicit declaration of function 'vm_flags_set' [-Werror=implicit-function-declaration]
 891 | vm_flags_set(vma, VM_MIXEDMAP);
      | ^~~~~~
drivers/rknpu/rknpu_gem.c: In function 'rknpu_gem_mmap_buffer':
drivers/rknpu/rknpu_gem.c:988:2: error: implicit declaration of function 'vm_flags_clear' [-Werror=implicit-function-declaration]
 988 | vm_flags_clear(vma, VM_PFNMAP);
      | ^~~~~~
cc1: all warnings being treated as errors
make[2]: *** [scripts/Makefile.build:273: drivers/rknpu/rknpu_gem.o] Error 1
make[1]: *** [scripts/Makefile.build:516: drivers/rknpu] Error 2
make[1]: *** Waiting for unfinished jobs....
AR      drivers/iio/built-in.a
make: *** [Makefile:1929: drivers] Error 2
MAKE KERNEL IMAGE FAILED.
```

Figure 2-1: Example of Kernel Compilation Error

Then need to modify the kernel header file kernel/include/linux/mm.h by adding the following code:

```
static inline void vm_flags_set(struct vm_area_struct *vma,
                               vm_flags_t flags){
    vma->vm_flags |= flags;
}

static inline void vm_flags_clear(struct vm_area_struct *vma,
                                  vm_flags_t flags){
    vma->vm_flags &= ~flags;
}
```

Rockchip

3 RKLLM Instruction Guide

3.1 Model Conversion

RKLLM-Toolkit provides model conversion and quantization functionalities. As one of the core features of RKLLM-Toolkit, it allows users to convert LLMs in Hugging Face format or GGUF format into RKLLM models, enabling the deployment and execution of RKLLM models on Rockchip NPU. This section will focus on the specific implementation of model conversion by RKLLM-Toolkit for LLMs, serving as a reference for users.

3.1.1 RKLLM Initialization

In this section, users need to initialize the RKLLM object, which is the first step in the entire workflow. In the example code, use the RKLLM() constructor function to initialize the RKLLM object.

```
rkllm = RKLLM()
```

3.1.2 Loading Models

After initializing RKLLM, users need to call the rkllm.load_huggingface() function to pass the specific path of the model. RKLLM-Toolkit will then successfully load the large language model in Hugging Face format or GGUF format based on the corresponding path, enabling subsequent conversion and quantization operations. The specific function definition is as follows:

Table 3-1 Interface Specification for the load_huggingface Function

Fuctionom	load_huggingface
Introduction	Used to load open-source LLMs in Hugging Face format.
Parameters	model: The path where the LLM files are stored, used for loading the model for subsequent conversion and quantization.
Returns	0 indicates successful model loading; -1 indicates model loading failure.

The example code is as follows:

```
ret = rkllm.load_huggingface(  
    model = './huggingface_model_dir',  
    model_lora = './huggingface_lora_model_dir')  
if ret != 0:  
    print('Load model failed!')
```

Table 3-2 Interface Specification for the load_gguf Function

Fuctionm	load_gguf
Introduction	Used to load open-source large language models in GGUF format, supporting the numerical types q4_0 and fp16. GGUF-format Lora models can also be loaded and converted to RKLLM models through this interface.
Parameters	model: The path where the LLM files are stored, used for loading the model for subsequent conversion and quantization. model_lora: The path of the Lora weights; when in converting, the model must be set to the corresponding base model path.
Returns	0 indicates successful model loading; -1 indicates model loading failure.

The example code is as follows:

```
ret = rkllm.load_gguf(model = './model-Q4_0.gguf')  
if ret != 0:  
    print('Load model failed!')
```

3.1.3 Construction for RKLLM Model

After loading the original model through the `rkllm.load_huggingface()` function, the next step is to build the RKLLM model using the `rkllm.build()` function. During model conversation, users can choose whether to perform quantization, which helps reduce the model size and improve inference performance on Rockchip NPU. The specific definition of the `rkllm.build()` function is as follows:

Table 3-3 Interface Specification for the build Function

Fuctionm	build
Introduction	Used to construct the RKLLM model and define specific quantization operations during the conversion process.

Parameters	<p><i>do_quantization:</i> This parameter controls whether to perform quantization operations on the model, and it is recommended to set it to True.</p> <p><i>optimization_level:</i> This parameter is used to set whether to perform quantization precision optimization. The available settings are {0, 1}, where 0 means no optimization and 1 means precision optimization. Precision optimization may cause a decrease in model inference performance.</p> <p><i>quantized_dtype:</i> This parameter is used to set the specific type of quantization. Currently supported types include "w4a16", "w4a16_g32", "w4a16_g64", "w4a16_g128", "w8a8", "w8a8_g128", "w8a8_g256", "w8a8_g512". "w4a16" means 4-bit quantization for weights and no quantization for activations. "w4a16_g64" means 4-bit grouped quantization for weights (group size=64) and no quantization for activations. "w8a8" means 8-bit quantization for both weights and activations. "w8a8_g128" means 8-bit grouped quantization for both weights and activations (group size=128). The rk3576 platform supports five quantization types: "w4a16", "w4a16_g32", "w4a16_g64", "w4a16_g128", and "w8a8". The rk3588 platform supports four quantization types: "w8a8", "w8a8_g128", "w8a8_g256", and "w8a8_g512". For GGUF models, the quantization type corresponding to q4_0 is "w4a16_g32". Note: The group size must divide the output dimension of the linear layer; otherwise, quantization will fail!</p> <p><i>quantized_algorithm:</i> Quantization accuracy optimization algorithm, with selectable options being "normal" or "gdq". All quantization types can choose "normal", but the "gdq" algorithm only supports "w4a16" and grouped "w4a16" quantization, and it requires high computational power, with GPU acceleration being necessary.</p> <p><i>num_npu_core:</i> The number of NPU cores to be used for model inference. For "rk3576", the options are [1, 2], and for "rk3588", the options are [1, 2, 3].</p>
------------	--

	<p>extra_qparams: Using the "gdq" algorithm generates a gdq.qparams quantized weight cache file. This parameter can be set to the gdq.qparams path to rebegin the model export.</p> <p>dataset: The dataset used for quantization calibration, formatted as JSON. An example of the content is as follows, where input is the question (with a prompt), and target is the answer. Multiple data entries are saved as dictionaries in a list;</p> <p>hybrid_rate: The hybrid quantization rate ($\in [0,1)$). When the quantization type is "w4a16"/"w8a8", the model will mix "w4a16_g"/"w8a8_g" types at the specified rate to improve accuracy. When the quantization type is grouped "w4a16_g"/"w8a8_g", the model will mix "w4a16"/"w8a8" types at the specified rate to improve inference performance. When the hybrid_rate value is 0, no hybrid quantization will be performed.</p> <p>target_platform: The hardware platform for running the model, with selectable options including "rk3576" or "rk3588".</p>
Returns	<p>0 indicates successful model conversion and quantization;</p> <p>-1 indicates model conversion failure.</p>

The example code is as follows:

```
ret = rkllm.build(
    do_quantization=True,
    optimization_level=1,
    quantized_dtype='w8a8',
    quantized_algorithm="normal",
    num_npu_core=3,
    extra_qparams=None,
    dataset="quant_data.json",
    hybrid_rate=0,
    target_platform='rk3588')
if ret != 0:
    print('Build model failed!')
```

3.1.4 Exporting Model

After constructing the RKLLM model using the rkllm.build() function, users can save the RKNN model as a .rkllm file for subsequent model deployment using the rkllm.export_rkllm() function. The

specific parameter definition of the `rkllm.export_rkllm()` function is as follows:

Table 3-4 Interface Specification for the `export_rkllm` Function

Fuctionom	<code>export_rkllm</code>
Introduction	Used to save the converted and quantized RKLLM model for subsequent inference.
Parameters	<i>export_path</i> : The save path for exporting the RKLLM model file. Lora models will automatically be saved with an <code>_lora</code> suffix in the RKLLM model filename.
Returns	0 indicates successful export and saving of the model; -1 indicates model export failure.

The example code is as follows:

```
ret = rkllm.export_rkllm(export_path = './model.rkllm')
if ret != 0:
    print('Export model failed!')
```

3.1.5 GPTQ Model Conversion

The user can convert a floating-point model to a 4-bit model using the [AutoGPTQ](#) open-source quantization tool before converting it to an RKLLM model, in addition to using the quantization algorithms provided in the tools above. When using AutoGPTQ to quantize the floating-point model, the following parameters must be configured:

```
bits=4
sym=true
group_size=32/64/128
desc_act=false
```

Here' s an example code snippet for converting a GPTQ model in Hugging Face format to an RKLLM model:

```

modelpath = '/path/to/Model-Instruct-GPTQ-Int4'
llm = RKLLM()

ret = llm.load_huggingface(model=modelpath, model_lora = None,
device='cuda')
if ret != 0:
    print('Load model failed!')
    exit(ret)

# Build model
dataset = None
qparams = None
target_platform = "RK3576"
optimization_level = 1
quantized_dtype = "w4a16_g32" #w4a16_g64 or w4a16_g128
quantized_algorithm = "normal"
num_npu_core = 2

ret = llm.build(do_quantization=True,
optimization_level=optimization_level,
quantized_dtype=quantized_dtype,
                quantized_algorithm=quantized_algorithm,
target_platform=target_platform, num_npu_core=num_npu_core,
extra_qparams=qparams, dataset=dataset)
if ret != 0:
    print('Build model failed!')
    exit(ret)

# Export rkllm model
ret =
llm.export_rkllm(f"./{os.path.basename(modelpath)}_{quantized_dtype}_{
target_platform}.rkllm")
if ret != 0:
    print('Export model failed!')
    exit(ret)

```

3.1.6 Update For Old Version Model

The users can update a model from version 1.0.2 to version 1.1 by using the `rkllm.update_rkllm()` function. When updating the model, there is no need to perform the model loading and building steps mentioned earlier. Instead, the update can be done directly via this interface, and after the update, parameters such as model quantization type remain unchanged.

Here is a description of the function and its parameters:

Table 3-5 Interface Specification for the update_rkllm Function

Fuctionom	update_rkllm
Introduction	Update the model from version 1.0.2 to version 1.1.
Parameters	model: The path to the RKLLM model of version 1.0.2.
Returns	0 indicates that the model update was successful; -1 indicates that the model update failed.

The example code is as follows:

```
ret = llm.update_rkllm(model = "./model_1.0.2version.rkllm")
if ret != 0:
    print('Load model failed!')
    exit(ret)
```

3.1.7 Simulation Accuracy Evaluation

After the user builds the RKLLM model through the rkllm.build() function, they can perform simulation accuracy evaluation on the PC using the rkllm.get_logits() function. The specific parameter definitions for the rkllm.get_logits() function are as follows:

Table 3-6 Interface Specification for the get_logits Function

Fuctionom	get_logits
Introduction	Used for simulation accuracy evaluation on the PC.
Parameters	inputs: The simulation input format is the same as Hugging Face model inference. An example is as follows: {"input_ids": "", "top_k": 1, ...}
Returns	The logits values inferred by the model.

Example code for using this function to perform PPL (perplexity) testing on the Wikitext dataset is as follows:

```
def eval_wikitext(llm):
    seqlen = 512
    tokenizer = AutoTokenizer.from_pretrained(
        modelpath,
        trust_remote_code=True
    )
    #Dataset download link:
    #https://huggingface.co/datasets/Salesforce/wikitext/tree/main/wikitext-2-raw-v1
    testenc = load_dataset("parquet", data_files='./wikitext/wikitext-2-raw-1/test-00000-of-00001.parquet', split='train')
    testenc = tokenizer(
        "\n\n".join(testenc['text']),
        return_tensors="pt").input_ids
    nsamples = testenc.numel() // seqlen
    nlls = []
    for i in tqdm(range(nsamples), desc="eval_wikitext: "):
        batch = testenc[:, (i * seqlen): ((i + 1) * seqlen)]
        inputs = {"input_ids": batch}
        lm_logits = llm.get_logits(inputs)
        if lm_logits is None:
            print("get logits failed!")
            return
        shift_logits = lm_logits[:, :-1, :]
        shift_labels = batch[:, 1:].to(lm_logits.device)
        loss_fct = nn.CrossEntropyLoss().to(lm_logits.device)
        loss = loss_fct(shift_logits.view(-1, shift_logits.size(-1)),
            shift_labels.view(-1))
        neg_log_likelihood = loss.float() * seqlen
        nlls.append(neg_log_likelihood)
    ppl = torch.exp(torch.stack(nlls).sum() / (nsamples * seqlen))
    print(f'wikitext-2-raw-1-test ppl: {round(ppl.item(), 2)}')
```

3.1.8 Simulated Model Inference

After the user builds the RKLLM model using the `rkllm.build()` function, they can perform simulated inference on the PC using the `rkllm.chat_model()` function. The specific parameter definitions for the `rkllm.chat_model()` function are as follows:

Table 3-7 Interface Specification for the `chat_model` Function

Fuctionom	<code>chat_model</code>
Introduction	Used for simulate model inference on the PC.
Parameters	<p>messages: The text input, which needs to include the appropriate prompts.</p> <p>args: Inference configuration parameters, such as sampling parameters like <code>top_k</code>.</p>
Returns	The logits values inferred by the model.

The example code is as follows:

```
args = {
    "max_length":128,
    "top_k":1,
    "temperature":0.8,
    "do_sample":True,
    "repetition_penalty":1.1
}

mesg = "Human: How's the weather today?\nAssistant:"
print(llm.chat_model(mesg, args))
```

The above operations cover all steps of model conversion and quantization in the RKLLM-Toolkit.

Depending on different requirements and application scenarios, users can choose different configuration options and quantization methods for customised settings, which facilitates subsequent deployment.

3.2 Inference Implementation in Board-side

This chapter introduces the usage of the general API interface functions. Users can refer to the content of this chapter to construct C++ code and implement inference of RKLLM models on the board to obtain inference results. The RKLLM board-side inference implementation is as follows:

- 1) Define the callback function `callback()`.
- 2) Define the RKLLM model parameter structure `RKLLMParam`.
- 3) Initialize the RKLLM model with `rkllm_init()`.
- 4) Perform model inference with `rkllm_run()`.
- 5) Process the real-time inference results returned by the callback function `callback()`.
- 6) Destroy the RKLLM model and release resources with `rkllm_destroy()`.

In the subsequent parts of this chapter, the document will provide detailed explanations of each step in the process and provide detailed explanations of the functions involved.

3.2.1 Define Callback Function

The callback function is used to receive real-time output results from the RKLLM model. It is bound during the initialization of RKLLM and continuously outputs results to the callback function during the RKLLM model inference process, returning only one token each time.

Here is an example that callback function prints the output results in real-time to the terminal:

```
void callback(RKLLMResult* result, void* userdata, LLMCallState state)
{
    if(state == LLM_RUN_NORMAL){
        printf("%s", result->text);
        for (int i=0; i<result->num; i++) {
            printf("token_id: %d logprob: %f", result->tokens[i].id,
                result->tokens[i].logprob);
        }
    }
    if (state == LLM_RUN_FINISH) {
        printf("finish\n");
    } else if (state == LLM_RUN_ERROR){
        printf("\run error\n");
    }
}
```

1) LLMCallState is a status flag, and its specific definition is as follows:

Table 3-8 Explanation of LLMCallState Status Flags

Definition	LLMCallState
Introduction	Used to indicate the current running state of RKLLM.
Enumeration Values	<p>LLM_RUN_NORMAL: indicates that the RKLLM model is currently inferencing;</p> <p>LLM_RUN_FINISH: indicates that the RKLLM model has completed inference;</p> <p>LLM_RUN_WAITING: indicates that the currently decoded character from RKLLM is not a complete UTF-8 encoding and needs to be concatenated with the next decoding result;</p> <p>LLM_RUN_ERROR: indicates that an error has occurred during inference;</p>

During the design process of the callback function, users can set different post-processing behaviors based on the different states of LLMCallState.

2) RKLLMResult is a return value structure, and its specific definition is as follows:

Table 3-9 Explanation of RKLLMResult Structure

Definition	RKLLMResult
Introduction	Used to return the current inference-generated result.
Struct Fields	<p>const char* text: indicates the text content generated by the current inference;</p> <p>int32_t token_id: indicates the token_id generated by the current inference;</p>

During the design process of the callback function, users can set different post-processing

behaviors based on the values in RKLLMResult.

3.2.2 Define RKLLMParam

The RKLLMParam structure is used to describe and define the detailed information of RKLLM. The specific definition is as follows:

Table 3-10 Explanation of RKLLMParam Structure

Definition	RKLLMParam
Introduction	Used to define the detailed parameters of the RKLLM model.
Struct Fields	<p><i>const char* model_path</i>: the path to the RKLLM model file;</p> <p><i>int32_t num_npu_core</i>: the number of NPU cores used during model inference; The "rk3576" platform has configurable range [1, 2]; while the "rk3588" is [1, 3];</p> <p><i>bool use_gpu</i>: whether to use GPU for prefill acceleration, default option is false;</p> <p><i>int32_t max_context_len</i>: the maximum context length during inference;</p> <p><i>int32_t max_new_tokens</i>: the maximum number of generated tokens in inferencing;</p> <p><i>int32_t top_k</i>: top-k sampling is a text generation method that selects the next token only from the top-k tokens with the highest probabilities predicted by the model. This method helps reduce the risk of generating low-probability or meaningless tokens. A higher top-k value (e.g., 100) will consider more token choices, resulting in more diverse text generation, while a lower value (e.g., 10) will focus on the most probable tokens, generating more conservative text. The default value is 40;</p> <p><i>float top_p</i>: top-p sampling, also known as nucleus sampling, is another text generation method that selects the next token from a group of tokens with cumulative probabilities of at least p. This method balances diversity and quality by considering the probabilities of tokens and the number of sampled tokens. A higher top-p value (e.g., 0.95) results in more diverse text generation, while a lower value (e.g., 0.5) generates more focused and conservative text. The default value is 0.9;</p>

	<p><i>float temperature:</i> a hyperparameter that controls the randomness of generated text by adjusting the probability distribution of output tokens. A higher temperature (e.g., 1.5) makes the output more random and creative. When the temperature is high, the model considers more options with lower probabilities when selecting the next token, resulting in more diverse and unexpected outputs. A lower temperature (e.g., 0.5) makes the output more focused and conservative. Lower temperatures mean that the model is more likely to choose high-probability tokens, resulting in more consistent and predictable outputs. In the extreme case of a temperature of 0, the model always chooses the most probable next token, resulting in identical outputs every time. To balance randomness and determinism and ensure that the output is neither overly uniform and predictable nor overly random and chaotic, the default value is 0.8;</p> <p><i>float repeat_penalty:</i> controls the occurrence of token sequence repetitions in the generated text, helping to prevent the model from generating repetitive or monotonous text. A higher value (e.g., 1.5) imposes a stronger penalty on repetitions, while a lower value (e.g., 0.9) is more lenient. The default value is 1.1;</p> <p><i>float frequency_penalty:</i> a factor for penalizing word/phrase repetition, reducing the probability of using words/phrases with higher frequencies overall and increasing the likelihood of using those with lower frequencies. This may lead to more diversified generated text, but could also result in text that is difficult to understand or not as expected. The range is [-2.0, 2.0], with a default value of 0;</p> <p><i>int32_t mirostat:</i> an algorithm actively maintaining the quality of generated text within the expected range during the text generation process. It aims to find a balance between coherence and diversity, avoiding low-quality output caused by excessive repetition (boredom trap) or incoherence (confusion trap). The values space is {0, 1, 2}, where 0 indicates not activating the algorithm, 1 indicates using the mirostat algorithm, and 2 indicates using the mirostat 2.0 algorithm;</p>
--	---

	<p><i>float mirostat_tau</i>: an option setting the target entropy for mirostat, representing the expected perplexity value of the generated text. Adjusting the target entropy allows to control the balance between coherence and diversity in the generated text. Lower values will result in more concentrated and coherent text, while higher values will lead to more diversified text, possibly with lower coherence. The default value is 5.0;</p> <p><i>float mirostat_eta</i>: an option setting the learning rate for mirostat, which influences the algorithm's responsiveness to feedback on generated text. A lower learning rate will result in slower adjustment, while a higher learning rate will make the algorithm more sensitive. The default value is 0.1;</p> <p><i>bool skip_special_token</i>: whether to skip special tokens and not output them, such as the end-of-sequence token <EOS>.</p> <p><i>bool is_async</i>: whether to use asynchronous mode.</p> <p><i>const char img_start*</i>: option to set the start marker for multimodal input image encoding, which needs to be configured in multimodal input mode.</p> <p><i>const char img_end*</i>: option to set the end marker for multimodal input image encoding, which needs to be configured in multimodal input mode.</p> <p><i>const char img_content*</i>: option to set the content marker for multimodal input image encoding, which needs to be configured in multimodal input mode.</p> <p><i>RKLLMExtendParam extend_param</i>: the special parameters for controlling inference.</p>
--	--

Table 3-11 Explanation of RKLLMExtendParam Structure

Definition	RKLLMExtendParam
Introduction	The special parameters for controlling inference.
Struct Fields	<p><i>int32_t base_domain_id</i>: controls from which domain the RKLLM model starts initialization, default is 0.</p>

In actual code construction, RKLLMParam needs to call the rkllm_createDefaultParam() function to

initialize its definition, and set the corresponding model parameters according to requirements. Sample code is as follows:

```
RKLLMParam param = rkllm_createDefaultParam();
param.model_path = "model.rkllm";
param.top_k = 1;
param.max_new_tokens = 256;
param.max_context_len = 512;
```

3.2.3 Definite Input Struct

To accommodate different input data types, the RKLLMInput input struct is defined, which currently accepts four types of input: text, image and text, token IDs, and encoded vectors. The specific definition is as follows:

Table 3-12 Explanation of RKLLMInput Structure

Definition	RKLLMInput
Introduction	Used to receive different forms of input data.
Struct Fields	<p><i>RKLLMInputType input_type</i>: Input mode;</p> <p><i>union</i>: Used to store different input data types, specifically including the following forms:</p> <p><i>const char prompt_input*</i>: Text prompt input, used to pass natural language text;</p> <p><i>RKLLMEmbedInput embed_input</i>: Embedding vector input, representing processed feature vectors;</p> <p><i>RKLLMTokenInput token_input</i>: Token input, used to pass the tokenized token sequence;</p> <p><i>RKLLMMultiModelInput multimodal_input</i>: Multimodal input, which can pass multimodal data, such as combined input of images and text;</p>

Table 3-13 Explanation of RKLLMInputType Structure

Definition	RKLLMInputType
Introduction	Used to represent the type of input data.
Enumeration	<i>RKLLM_INPUT_PROMPT</i> : Indicates that the input data is plain text.

Values	<p><i>RKLLM_INPUT_TOKEN</i>: Indicates that the input data is token IDs.</p> <p><i>RKLLM_INPUT_EMBED</i>: Indicates that the input data is encoded vectors.</p> <p><i>RKLLM_INPUT_MULTIMODAL</i>: Indicates that the input data consists of images and text.</p>
--------	---

When the input data is plain text, it can be directly input using `input_data`. When the input data consists of token IDs, encoded vectors, or images and text, the `RKLLMInput` must be used in conjunction with the `RKLLMTokenInput`, `RKLLMEmbedInput`, and `RKLLMMultiModelInput` structures. The specific introduction is as follows:

1) `RKLLMTokenInput` is the input struct that receives token IDs. The specific definition is as follows:

Table 3-14 Explanation of `RKLLMTokenInput` Structure

Definition	<code>RKLLMTokenInput</code>
Introduction	Used to receive <code>token_id</code> data.
Struct Fields	<p><i>int32_t input_ids*</i>: Memory pointer for the input token IDs.</p> <p><i>size_t n_tokens</i>: The number of tokens in the input data.</p>

2) `RKLLMEmbedInput` is the input struct that receives encoded vectors. The specific definition is as follows:

Table 3-15 Explanation of `RKLLMEmbedInput` Structure

Definition	<code>RKLLMEmbedInput</code>
Introduction	Used to receive embedding data.
Struct Fields	<p><i>float embed*</i>: Memory pointer for the input token embeddings.</p> <p><i>size_t n_tokens</i>: The number of tokens in the input data.</p>

3) `RKLLMMultiModelInput` is the input struct that receives images and text. The specific definition is as follows:

Table 3-16 Explanation of `RKLLMMultiModelInput` Structure

Definition	<code>RKLLMMultiModelInput</code>
------------	-----------------------------------

Introduction	Used to receive images and text data.
Struct Fields	<p><i>char prompt*</i>: Memory pointer for the input text.</p> <p><i>float image_embed*</i>: Memory pointer for the input image embeddings.</p> <p><i>size_t n_image_tokens</i>: The number of tokens for the input image embeddings.</p>

Here is an example of pure text input code:

```
#define PROMPT_TEXT_PREFIX "<|im_start|>system You are a helpful assistant. <|im_end|> <|im_start|>user"
#define PROMPT_TEXT_POSTFIX "<|im_end|><|im_start|>assistant"

string input_str = "把这句话翻译成英文: RK3588 是新一代高端处理器, 具有高算力、低功耗、超强多媒体、丰富数据接口等特点";
input_str = PROMPT_TEXT_PREFIX + input_str + PROMPT_TEXT_POSTFIX;

RKLLMInput rkllm_input;
rkllm_input.input_data = (char*)input_str.c_str();
rkllm_input.input_type = RKLLM_INPUT_PROMPT;

RKLLMInferParam rkllm_infer_params;
memset(&rkllm_infer_params, 0, sizeof(RKLLMInferParam));
rkllm_infer_params.mode = RKLLM_INFER_GENERATE;
```

An example code for image and text multimodal input is as follows. Note that the prompt for multimodal input needs to include the <image> placeholder to indicate where the image encoding should be inserted:

```
#define PROMPT_TEXT_PREFIX "<|im_start|>system You are a helpful assistant. <|im_end|> <|im_start|>user"
#define PROMPT_TEXT_POSTFIX "<|im_end|><|im_start|>assistant"

RKLLMInput rkllm_input;
string input_str = "<image>Please describe the image shortly.";
input_str = PROMPT_TEXT_PREFIX + input_str + PROMPT_TEXT_POSTFIX;
rkllm_input.multimodal_input.prompt = (char*)input_str.c_str();

rkllm_input.multimodal_input.n_image_tokens = 256;
int rkllm_input_len = multimodal_input.n_image_tokens * 3072;
rkllm_input.multimodal_input.image_embed = (float *)malloc(rkllm_input_len * sizeof(float));
FILE *file;
file = fopen("models/image_embed.bin", "rb");
fread(rkllm_input.multimodal_input.image_embed, sizeof(float), rkllm_input_len, file);
fclose(file);

rkllm_input.input_type = RKLLM_INPUT_MULTIMODAL;
RKLLMInferParam rkllm_infer_params;
memset(&rkllm_infer_params, 0, sizeof(RKLLMInferParam));
rkllm_infer_params.mode = RKLLM_INFER_GENERATE;
```

RKLLM supports different inference modes and defines the RKLLMInferParam structure. It currently

supports joint inference with preloaded LoRA models during the inference process, or saving a Prompt Cache for subsequent inference acceleration. The specific definition is as follows:

Table 3-17 Explanation of RKLLMIInferParam Structure

Definition	RKLLMIInferParam
Introduction	Used to define different inference modes.
Struct Fields	<p><i>RKLLMIInferMode mode</i>: Inference mode, currently, only the standard inference mode RKLLM_INFER_GENERATE is supported.</p> <p><i>RKLLMLoraParam lora_params*</i>: Parameter configuration for the LoRA used during inference, used to select which LoRA to infer when multiple LoRAs are loaded. Set to NULL if LoRA is not needed.</p> <p><i>RKLLMPromptCacheParam prompt_cache_params*</i>: Parameter configuration for using the Prompt Cache during inference. Set to NULL if Prompt Cache generation is not needed.</p>

Table 3-18 Explanation of RKLLMLoraParam Structure

Definition	RKLLMLoraParam
Introduction	Used to define the parameters for using LoRA during inference.
Struct Fields	<i>const char lora_adapter_name*</i> : The name of the LoRA used during inference.

Table 3-19 Explanation of RKLLMPromptCacheParam Structure

Definition	RKLLMPromptCacheParam
Introduction	Used to define the parameters for using Prompt Cache during inference.
Struct Fields	<p><i>int save_prompt_cache</i>: Indicates whether to save the Prompt Cache during inference. 1 means it is required, and 0 means it is not.</p> <p><i>const char prompt_cache_path*</i>: Path to save the Prompt Cache. If not set, it defaults to “./prompt_cache.bin”.</p>

Here is an example of using InferParam for inference:

```
// 1. Initialize and set LoRA parameters (if needed)
RKLLMLoraParam lora_params;
// Specify the LoRA model name
lora_params.lora_adapter_name = "test";
// 2. Initialize and Set Prompt Cache Parameters(if needed)
RKLLMPromptCacheParam prompt_cache_params;
// Enable saving Prompt Cache
prompt_cache_params.save_prompt_cache = true;
// Specify cache file path
prompt_cache_params.prompt_cache_path = "./prompt_cache.bin";
rkllm_infer_params.mode = RKLLM_INFER_GENERATE;
rkllm_infer_params.lora_params = &lora_params;
rkllm_infer_params.prompt_cache_params = &prompt_cache_params;
```

3.2.4 Initialize Model

Before initializing the model, it is necessary to define the LLMHandle handle in advance. This handle is used for the initialization, inference, and resource release processes of the model. It's important to note that only by unifying the LLMHandle handle object across these three processes can the inference process of the model be completed correctly.

Prior to model inference, users need to complete the model initialization through the rkllm_init() function. The specific function definition is as follows:

Table 3-20 Interface Specification for the rkllm_init Function

Fuctionom	rkllm_init
Introduction	Used to initialize the specific parameters and inference settings for RKLLM model.
Parameters	<p>LLMHandle* handle: register model to the corresponding handle for subsequent inference and release calls;</p> <p>RKLLMParam* param: the parameter structure defined for the model;</p> <p>LLMResultCallback callback: callback function used to receive and process real-time outputs from the model;</p>
Returns	<p>0 indicates that the initialization process is normal;</p> <p>-1 indicates initialization failure;</p>

The example code is as follows:

```
LLMHandle llmHandle = nullptr;
rkllm_init(&llmHandle, &param, callback);
```

3.2.5 Inference Model

After completing the initialization process of the RKLLM model, users can perform model inference using the `rkllm_run()` function. Real-time inference results can be processed using the callback function predefined during initialization. The specific function definition of `rkllm_run()` is as follows:

Table 3-21 Interface Specification for the `rkllm_run` Function

Fuctionom	<code>rkllm_run</code>
Introduction	Used to performing result inference using the initialized RKLLM model.
Parameters	<p><i>LLMHandle handle</i>: the target handle registered during model initialization.</p> <p><i>RKLLMInput rkllm_input*</i>: Input data for model inference. For details, see section 3.2.3 on input structure definition.</p> <p><i>RKLLMInferParam rkllm_infer_params*</i>: Parameter passing during the model inference process. For details, see section 3.2.3 on input structure definition.</p> <p><i>void* userdata</i>: the user-defined function pointer, typically set to NULL by default.</p>
Returns	<p>0 indicates that the model inference runs normally;</p> <p>-1 indicates a failure in calling the model inference;</p>

The example code is as follows:

```
#define PROMPT_TEXT_PREFIX "<|im_start|>system You are a helpful assistant. <|im_end|> <|im_start|>user"
#define PROMPT_TEXT_POSTFIX "<|im_end|><|im_start|>assistant"

// Predefined text values for the prompt before and after
string input_str = "把这句话翻译成英文：RK3588 是新一代高端处理器，具有高算力、低功耗、超强多媒体、丰富数据接口等特点";
input_str = PROMPT_TEXT_PREFIX + input_str + PROMPT_TEXT_POSTFIX;

// Define the input prompt and complete the concatenation
RKLLMInferParam rkllm_infer_params;
memset(&rkllm_infer_params, 0, sizeof(RKLLMInferParam));
rkllm_infer_params.mode = RKLLM_INFER_GENERATE;
// 1. Initialize and set LoRA parameters (if needed)
RKLLMLoraParam lora_params;
lora_params.lora_adapter_name = "test";
// 2. Initialize and Set Prompt Cache Parameters(if needed)
RKLLMPromptCacheParam prompt_cache_params;
prompt_cache_params.save_prompt_cache = true;
prompt_cache_params.prompt_cache_path = "./prompt_cache.bin";
rkllm_infer_params.mode = RKLLM_INFER_GENERATE;
```

```
// rkllm_infer_params.lora_params = &lora_params;
// rkllm_infer_params.prompt_cache_params = &prompt_cache_params;
rkllm_infer_params.lora_params = NULL;
rkllm_infer_params.prompt_cache_params = NULL;

RKLLMInput rkllm_input;
rkllm_input.input_type = RKLLM_INPUT_PROMPT;
rkllm_input.prompt_input = (char *)text.c_str();

rkllm_run(llmHandle, &rkllm_input, &rkllm_infer_params, NULL);
```

3.2.6 Interrupt Model Inference

During model inference, users can call the `rkllm_abort()` function to interrupt the inference process.

The specific function definition is as follows:

Table 3-22 Interface Specification for the `rkllm_abort` Function

Fuctionm	<code>rkllm_abort</code>
Introduction	Used to interrupt the RKLLM model inference process.
Parameters	<i>LLMHandle handle</i> : the target handle registered during model initialization;
Returns	0 indicates successful interruption of the model; -1 indicates a failure to interrupt the model.

The example code is as follows:

```
// llmHandle is the the target handle registered
rkllm_abort(llmHandle);
```

3.2.7 Release Model

After completing all model inference calls, users need to call the `rkllm_destroy()` function to destroy the RKLLM model and release the CPU, GPU, and NPU computing resources allocated, for use by other processes or models. The specific function definition is as follows:

Table 3-23 Interface Specification for the `rkllm_destory` Function

Fuctionm	<code>rkllm_destroy</code>
Introduction	Used to destroy the RKLLM model and release all computing resources.
Parameters	<i>LLMHandle handle</i> : the target handle registered during model initialization;

Returns	<p>0 indicates successful destruction and release of the RKLLM model;</p> <p>-1 indicates a failure in releasing the model.</p>
---------	---

The example code is as follows:

```
// llmHandle is the the target handle registered
rkllm_destroy(llmHandle);
```

3.2.8 Load LoRA Model

RKLLM supports running LoRA models simultaneously with the base model during inference. Before invoking the rkllm_run interface, you can load a LoRA model via the rkllm_load_lora interface. RKLLM allows loading multiple LoRA models; each call to rkllm_load_lora loads one LoRA model. The specific function definition is as follows:

Table 3-24 Interface Specification for the rkllm_load_lora Function

Fuctionom	rkllm_load_lora
Introduction	Used to load LoRA model for the base model.
Parameters	<p>LLMHandle handle: The target handle registered during model initialization. See section 3.2.4 on initializing the model.</p> <p>RKLLMLoraAdapter lora_adapter*: Parameter for loading the LoRA model.</p>
Returns	<p>0 indicates the LoRA model was successfully loaded.</p> <p>-1 indicates the model loading failed.</p>

Table 3-25 Explanation of RKLLMLoraAdapter Structure

Definition	RKLLMLoraAdapter
Introduction	Used to configure parameters when loading a LoRA model.
Struct Fields	<p>const char* lora_adapter_path: The path to the LoRA model to be loaded.</p> <p>const char* lora_adapter_name: The name of the LoRA model to be loaded, defined by the user, used to select the specified LoRA during inference.</p> <p>float scale: The degree to which the LoRA model adjusts the base model parameters during inference.</p>

Here is an example Code for Loading LoRA:

```
RKLLMLoraAdapter lora_adapter;
memset(&lora_adapter, 0, sizeof(RKLLMLoraAdapter));
lora_adapter.lora_adapter_path = "lora.rkllm";
lora_adapter.lora_adapter_name = "lora_name";
lora_adapter.scale = 1.0;
ret = rkllm_load_lora(llmHandle, &lora_adapter);
if (ret != 0) {
    printf("\nload lora failed\n");
}
```

3.2.9 Load Prompt Cache

During the model inference process, the Prefill phase typically consumes a large amount of computational resources and time, especially when the prompt is long. To speed up this process, RKLLM supports loading pre-generated Prompt Cache files. By reusing the content in the cache, the time spent on the Prefill phase can be significantly reduced, thereby improving overall inference efficiency. Before calling the `rkllm_run` interface for inference, ensure that the `prompt_cache_params` parameter is correctly configured. This step allows the model to generate the corresponding Prompt Cache file after inference is completed. When running inference for the first time, the system will automatically generate a Prompt Cache file, which contains the intermediate results required for the Prefill phase, for subsequent use. In later inference tasks, the previously generated Prompt Cache file can be loaded by calling the `rkllm_load_prompt_cache` interface.

The specific function definition is as follows:

Table 3-26 Interface Specification for the `rkllm_load_prompt_cache` Function

Fuctionom	<code>rkllm_load_prompt_cache</code>
Introduction	Used to load Prompt Cache file.
Parameters	<p><i>LLMHandle handle</i>: The target handle registered during model initialization (refer to section 3.2.4 Initialization Model).</p> <p><i>const char* prompt_cache_path</i>: The path to the Prompt Cache file to be loaded.</p>
Returns	<p>0 indicates the Prompt Cache file was successfully loaded.</p> <p>-1 indicates loading failed.</p>

Table 3-27 Interface Specification for the rkllm_release_prompt_cache Function

Fuctionom	rkllm_release_prompt_cache
Introduction	Used to release the Prompt Cache.
Parameters	LLMHandle handle: The target handle registered during model initialization (refer to section 3.2.4 Initialization Model).
Returns	0 indicates that the Prompt Cache model was successfully released. -1 indicates that the model release failed.

Note:

RKLLM will detect the parts of the input that are the same as those in the `prompt_cache` from the beginning. If your input format is fixed as `PROMPT_PREFIX + text + PROMPT_POSTFIX`, you can generate the Prompt Cache for just the `PROMPT_PREFIX` part. After loading, you can reuse this part of the result in subsequent inferences.

RKLLM supports generating multiple Prompt Cache files. When different Prompt Cache files are needed, you can simply load the corresponding file. If you need to switch to another Prompt Cache file or no longer need the loaded Prompt Cache, please explicitly call the `rkllm_release_prompt_cache` interface to release it.

Here is an example Code for Loading Prompt Cache:

```
// Initialize and set the Prompt Cache parameters, then call the run
// interface to generate the Prompt Cache file.

RKLLMPromptCacheParam prompt_cache_params;
// Whether to save the prompt cache
prompt_cache_params.save_prompt_cache = true;
// If you need to save the prompt cache, specify the absolute path of
// the cache file.
prompt_cache_params.prompt_cache_path = "/data/prompt_cache.bin";
rkllm_infer_params.prompt_cache_params = &prompt_cache_params;

rkllm_infer_params.mode = RKLLM_INFER_GENERATE;
rkllm_input.input_type = RKLLM_INPUT_PROMPT;
rkllm_input.prompt_input = (char *)prompt.c_str();
rkllm_run(llmHandle, &rkllm_input, &rkllm_infer_params, NULL);

// Load the prompt cache file to reduce prefill time.
rkllm_load_prompt_cache(llmHandle, "./prompt_cache.bin");
```

```
if (ret != 0) {  
    printf("\nload Prompt Cache failed\n");  
}  
  
rkllm_run(llmHandle, &rkllm_input, &rkllm_infer_params, NULL);
```

3.3 Board-side Inference Example

The directory, examples/DeepSeek-R1-Distill-Qwen-1.5B_Demo, includes the C++ project for the inference on the board-side, and which is synchronously updated with this documentation. Furthermore, compilation scripts are provided for users to facilitate the compilation of the project and the completion of the board-side inference of the RKLLM model.

3.3.1 Complete Code of Example Project

The complete C++ code example for inference calls is located in the example/DeepSeek-R1-Distill-Qwen-1.5B_Demo/deploy/src directory of the toolchain. The llm_demo.cpp file serves as an example of large language model inference. These examples include the full process, including model initialization, inference, handling outputs with callback functions, and releasing model resources. Users can refer to the relevant code to implement custom functionality.

3.3.2 Instructions for Example Project

Under the directory of examples/DeepSeek-R1-Distill-Qwen-1.5B_Demo/deploy, not only does it contain sample code for invoking the RKLLM model inference, but also includes compilation scripts build-android.sh and build-linux.sh. This section will provide a brief explanation of how to use the sample code, taking compiling executable files for Linux systems as an example using the build-linux.sh script:

Firstly, users need to prepare cross-compilation tools on their own, noting that the recommended version of the compilation tools in section 2.3 is 10.2 or above. Subsequently, before the compilation process, users should replace the path to the cross-compilation tools in build-linux.sh themselves:

```
# Set the path of the cross-compiler GCC_COMPILER_PATH=~/gcc-arm-10.2-  
2020.11-x86_64-aarch64-none-linux-gnu/bin/aarch64-none-linux-gnu
```

Subsequently, users can use build-linux.sh to initiate the compilation process. Upon completion of the compilation, users will obtain the corresponding llm_demo program, which will be installed in the

install/demo_linux_aarch64/llm_demo directory. Following this, the executable file, library folder, and the RKLLM model (previously converted and quantized using the RKLLM-Toolkit tool) should be pushed to the board-side.

```
adb push install/demo_linux_aarch64 /data
adb push /PC/path/to/your/rkllm/model /data/demo_linux_aarch64
```

After completing the above steps, users can enter the terminal interface of the board using adb, and navigate to the corresponding /data/demo_linux_aarch64 directory. Then, the inference of RKLLM model on the board can be invoked using the following command:

```
adb shell
cd /data/demo_linux_aarch64
export LD_LIBRARY_PATH=./lib
./llm_demo /path/to/your/rkllm/model 1024 2048
```

With the above operations, users can enter the example inference interface, interact with the board-side model for inference, and obtain real-time inference results from the RKLLM model.

3.3.3 Monitor inference performance

To monitor the inference performance of RKLLM on the board like the above figure, you can use the command:

```
export RKLLM_LOG_LEVEL=1
```

Stage	Total Time (ms)	Tokens	Time per Token (ms)	Tokens per Second
Prefill	307.22	11	27.93	35.81
Generate	2103.63	15	142.91	7.00

Figure 3-1 RKLLM inference performance logs on the hardware platform

This will display the number of tokens processed and the inference time for both the Prefill and Generate stages after each inference, as shown in the figure below. This information will help you evaluate the performance by providing detailed logging of how long each stage of the inference process takes.

3.4 Implementation of Board-side Server

After using RKLLM-Toolkit to convert the model and obtain the RKLLM model, users can deploy server-side services on Linux development boards. This involves setting up a server on a Linux device and

exposing network interfaces to everyone in the local area network. Subsequently, the RKLLM model can be accessed by other users via the specified address, thus facilitating efficient and concise interaction. This section will introduce two different server deployment implementations.

1) RLM-Server-Flask based on Flask: Users can achieve API access between the client and server using request requests. In the provided RKLLM-Server-Flask example, the send-receive structure is specially set to be the same as the OpenAI-API interface, facilitating quick replacement for users on existing development bases.

2) RKLLM-Server-Gradio based on Gradio: By reference to the provided example, users can rapidly construct a web server for visual interaction. Furthermore, the example illustrates the utilisation of the Gradio API interface, which enables users to undertake secondary development.

Both examples of server implementations mentioned above are located in the examples/rkllm_server_demo directory. This directory contains specific code for both implementations, one-click deployment scripts, and API interface calling examples. Users can choose different examples for reference and further development. The directory structure is as follows:

```
examples/rkllm_server_demo
├── rkllm_server          # Board-side Deployment Required Files
│   ├── lib              # RKLLM Runtime
│   ├── fix_freq_rk3576.sh # Script for RK3576 Fixed Frequency
│   ├── fix_freq_rk3588.sh # Script for RK3588 Fixed Frequency
│   ├── flask_server.py   # RKLLM-Server-Flask Example
│   └── gradio_server.py  # RKLLM-Server-Gradio Example
├── build_rkllm_server_flask.sh # One-click Deployment Script -Flask
├── build_rkllm_server_gradio.sh # One-click Deployment Script -Gradio
├── chat_api_flask.py       # API Interface Example -Flask
├── chat_api_gradio.py     # API Interface Example -Gradio
└── Readme.md
```

3.4.1 Deployment Example of RKLLM-Server-Flask

In the deployment example of RKLLM-Server-Flask, the main focus is on using the Flask framework to set up the server-side. On the client-side, data is transmitted using the request-response structure to achieve API access. In both the server and client implementations of RKLLM-Server-Flask, special consideration is given to the calling method of the OpenAI-API. The example code ensures consistency by setting up data structures identical to those used in the OpenAI-API. This allows users to seamlessly

migrate their services by simply replacing the access interface after deploying RKLLM-Server-Flask, leveraging their existing development foundations in OpenAI-API development.

According to the [OpenAI-API usage documentation](#), users send specific data structures to the server during the API call process. The main contents are as follows:

```
{
  "model": "No models available",
  "messages": [
    {
      "role": "system",
      "content": "You are a helpful assistant."
    },
    {
      "role": "user",
      "content": "Hello!"
    }
  ],
  "stream": false,
}
```

Among them, "model" and "stream" specify the specific model to be called and whether to initiate streaming inference transmission, while the "content" data in "messages" is the crucial user input.

As for the data returned by the server, the structure of the data output by the OpenAI-API varies depending on whether streaming inference transmission is selected. The data content returned under non-streaming inference settings is as follows:

```
{
  "id": "chatcmpl-123",
  "object": "chat.completion",
  "created": 1677652288,
  "model": "gpt-3.5-turbo-0125",
  "system_fingerprint": "fp_44709d6fcb",
  "choices": [{
    "index": 0,
    "message": {
      "role": "assistant",
      "content": "\n\nHello there, how may I assist you today?",
    },
    "logprobs": null,
    "finish_reason": "stop"
  }],
  "usage": {
    "prompt_tokens": 9,
    "completion_tokens": 12,
    "total_tokens": 21
  }
}
```

When streaming inference transmission is not selected, the most important part is the "messages" content under the "choices" section, which represents the inference results provided by the model.

In contrast, when streaming inference transmission is enabled, the server returns a Response object, which includes the output results of the model at different points in time during the streaming inference process. The data content at each moment is as follows:

```
{
  "id": "chatcmpl-123",
  "object": "chat.completion.chunk",
  "created": 1677652288,
  "model": "gpt-3.5-turbo-0125",
  "system_fingerprint": "fp_44709d6fcb",
  "choices": [{
    "index": 0,
    "delta": {
      "role": "assistant",
      "content": "\n\nHello there, how may I assist you today?", },
    "logprobs": null,
    "finish_reason": "stop"
  }]
}
```

After receiving the data from streaming transmission, users need to focus on the "delta" data section within the "choices" part. Additionally, when "finish_reason" is empty (None), it indicates that the model is still in the inference state, and the data has not been fully generated yet. It's only when "finish_reason" returns "stop" that the streaming inference is considered finished.

In the provided deployment example code and API access examples for RKLLM-Server-Flask, you can see identical definitions of the transmission data structure, ensuring the generality of the deployed RKLLM-Server-Flask. This facilitates quick replacement for users who have previously used OpenAI-API. In the subsequent sections of this chapter, we will separately introduce the one-click deployment script for the server, important settings for server deployment implementation, and the script design for client-side API access.

3.4.1.1 Server-side: One-click Deployment Script for RKLLM-Server-Flask

The one-click deployment script for RKLLM-Server-Flask is named `build_rkllm_server_flask.sh` and is located in the `rkllm_server_demo` directory. This script helps users quickly set up the RKLLM-Server-Flask server on a Linux development board. Before using this script, users should note the following:

- 1) Ensure that the development board is connected to the network via an Ethernet cable. Use the "ifconfig" command in the adb shell to determine the specific IP address of the development board. The

RKLLM-Server-Flask will then set up the server with this IP address within the local area network and accept client access.

2) Users should have successfully converted the RKLLM model beforehand. Before executing the one-click deployment script, ensure that the RKLLM model has been pushed to the Linux board.

Users can directly invoke the `build_rkllm_server_flask.sh` script from their PC (not on a development board) to quickly deploy the RKLLM-Server-Flask server on a Linux development board.

The specific usage of the one-click deployment script `build_rkllm_server_flask.sh` is as follows:

```
./build_rkllm_server_flask.sh
--workshop [RKLLM-Server Working Path]
--model_path [Absolute Path of Converted RKLLM Model on Board] --
platform [Target Platform: rk3588/rk3576]
[--lora_model_path [Lora Model Path]]
[--prompt_cache_path [Prompt Cache File Path]]
```

The `'workshop'` parameter specifies the subsequent working directory of RKLLM-Server-Flask on the board. The `'model_path'` parameter indicates the absolute path of the RKLLM model on the board, which was converted using RKLLM-Toolkit, and RKLLM-Server-Flask will read the model from this path during operation. The `'platform'` parameter specifies the current platform type, either `rk3588` or `rk3576`. The `'lora_model_path'` and `'prompt_cache_path'` parameters are optional and can be used to specify file paths when the user needs to load a Lora model or use the prompt feature.

The following is a simple example of how to use the one-click deployment script `'build_rkllm_server_flask.sh'`:

```
./build_rkllm_server_flask.sh
--workshop /path/to/workshop
--model_path /path/to/model.rkllm
--platform rk3588
```

After executing the above command, the one-click deployment script will perform the following steps:

- 1) Check the Linux environment on the board.
- 2) Automatically install the Flask library if not already installed.
- 3) Push the necessary files under `rkllm_server_demo/rkllm_server` to the board.
- 4) Index the RKLLM model in the preset working directory for RKLLM-Server-Flask.

Once you see the message "RKLLM Model has been initialized successfully!" in the terminal, it indicates that the RKLLM-Server-Flask example has been successfully launched.

```
=====init....=====
rkllm-runtime version: 1.0.2b9, rknpv driver version: 0.9.7, platform: RK3588
load prompt cache from '/data/cw/prompt_cache.bin'
loaded a prompt cache with prompt size of 27 tokens
RKLLM Model has been initialized successfully!
=====
* Serving Flask app 'flask_server'
* Debug mode: off
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on all addresses (0.0.0.0)
* Running on http://127.0.0.1:8080
* Running on http://172.16.10.79:8080
Press CTRL+C to quit
```

Figure 3-2 Successful deployment of RKLLM-Server-Flask in terminal

By referring to the specific code logic in `build_rkllm_server_flask.sh`, users can understand the detailed deployment process of the RKLLM-Server-Flask example. This enables users to customize the deployment implementation of their server more flexibly. It is important to emphasize that in step 3 of the one-click deployment script, the script automatically synchronizes the current version of RKLLM Runtime to `rkllm_server/lib/librkllmrt.so`. This ensures that `flask_server.py` calls the current version of `librkllmrt.so` during runtime.

3.4.1.2 Server-side: Introductions for RKLLM-Server-Flask Example

In this section, we will outline and introduce the implementation approach of the deployment example for RKLLM-Server-Flask, helping users understand the construction logic of the example code for potential secondary development.

The deployment example of RKLLM-Server-Flask primarily relies on the Flask library to achieve the basic implementation of the server. Additionally, for RKLLM model inference, the `ctypes` library in Python is chosen to directly call the RKLLM Runtime library.

In the overall implementation of `rkllm_server/flask_server.py`, in order to call `librkllmrt.so` via `ctypes`, it's necessary to define relevant structures in Python based on the header file `rkllm.h` corresponding to `librkllmrt.so` beforehand. After the Flask server receives the struct data sent by users, it calls relevant functions to perform inference with the RKLLM model. The specific code implementation of `flask_server.py` mainly consists of the following steps:

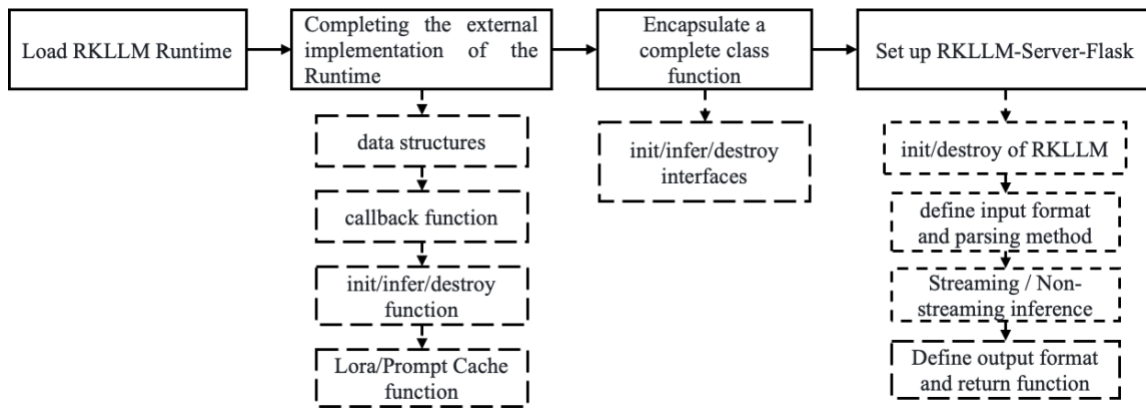


Figure 3-3 Overview of the RKLLM-Server-Flask Deployment Implementation Process

1) Load RKLLM Runtime: Set the path for the dynamic library and load the dynamic library `librkllmrt.so` using `ctypes` to achieve the analysis of the RKLLM Runtime.

2) Completing the external implementation of the Runtime: In the Python code, use `ctypes` to complete the external definitions of relevant implementations from the RKLLM Runtime header files, including definitions of data structures, callback functions, initialization functions, inference functions, destroy functions, and loading functions for Lora/Prompt Cache.

3) Encapsulate a complete class function: Based on the various Runtime data types and function interfaces implemented in step 2), encapsulate complete class functions that integrate RKLLM's initialization, inference, destroy, and other operations for easier subsequent calls.

4) Set up RKLLM-Server-Flask: Using the complete class functions encapsulated in step 3, build the Flask server, including loading the RKLLM model based on user-specified parameters at Flask startup, defining the format of user input and the method for parsing the input, differentiating between streaming/non-streaming inference call methods, defining the return format for inference output and the specific implementation of callback functions, and handling RKLLM release.

The specific implementations of the above modules form the main body of the code in `rkllm_server/flask_server.py`, thereby completing the deployment of the RKLLM-Server-Flask example. Users can modify the initialization definitions for the RKLLM model to implement different custom models. Additionally, users can refer to the RKLLM-Server-Flask deployment example for implementing their own custom server deployment.

3.4.1.3 Client-side: API Access Example

In `rkllm_server_demo/chat_api_flask.py`, an API access example for the aforementioned RKLLM-Server-Flask server is demonstrated. When developing custom functionalities, users only need to refer to this API access example and utilize corresponding send-receive structures for data wrapping and parsing. Since the send-receive data structures in this example follow the design of the OpenAI-API, users who have previously developed using the OpenAI-API only need to replace the corresponding network interfaces. This section will provide explanations for the important code segments:

1) Define the network address of RKLLM-Server-Flask. Users need to set the target address based on the specific IP of the Linux development board, the port number set by the Flask framework, and the function name for access.

```
server_url = 'http://172.16.10.166:8080/rkllm_chat'
```

2) Define the form of API access, with options for non-streaming transmission and streaming transmission, defaulting to non-streaming transmission.

```
is_streaming = True
```

3) Define the session object, using `requests.Session()` to configure the communication process between the API access interface and the server. Users can customize modifications according to their actual development needs.

```
session = requests.Session()
session.keep_alive = False # Close the connection pool
adapter = requests.adapters.HTTPAdapter(max_retries=5)
session.mount('https://', adapter)
session.mount('http://', adapter)
```

4) Define the structure used to wrap the data sent during API calls. User access to RKLLM-Server-Flask will be encapsulated within this structure.

```
# Prepare the data to be sent
# model: the model defined by the user when setting up RKLLM-Server,
# which has no effect here
# messages: the questions input by the user; supports adding multiple
# questions to messages
# stream: whether to enable streaming dialogue, same as the OpenAI
# interface
data = {
    "model": 'your_model_deploy_with_RKLLM_Server',
    "messages": [{"role": "user", "content": user_message}],
    "stream": is_streaming
}
```

5) Send a request to the RKLLM-Server-Flask server and retrieve the returned data.

```
responses = session.post(server_url, json=data, headers=headers,
stream=is_streaming, verify=False)
```

6) Define the parsing method for the returned data structure. Since RKLLM-Server-Flask follows the format of the returned struct from the OpenAI-API, parsing operations are required in actual usage.

```
# Parse the response
# Non-streaming transmission
if not is_streaming:
    if responses.status_code == 200:
        print("Q:", data["messages"][-1]["content"])
        print("A:", json.loads(responses.text)["choices"][-
1]["message"]["content"])
    else:
        print("Error:", responses.text)

# Streaming transmission
else:
    if responses.status_code == 200:
        print("Q:", data["messages"][-1]["content"])
        print("A:", end="")
        for line in responses.iter_lines():
if line:
            line = json.loads(line.decode('utf-8'))
            if line["choices"][-1]["finish_reason"] != "stop":
                print(line["choices"][-1]["delta"]["content"],
end="")

            sys.stdout.flush()
    else:
        print('Error:', responses.text)
```

The overall process from steps 1 to 6 represents the API access method for the RKLLM-Server-Flask server. Users can develop custom functionalities based on the example code provided above. It is important for users to verify the specific address of the RKLLM-Server-Flask, namely the IP address, port number, and the function interface that the server accepts input from. Additionally, when encountering specific requirements for send-receive structures, users can customize the required data structures on both the server and client sides to ensure the implementation of custom functionalities.

3.4.2 Deployment Example of RKLLM-Server-Gradio

Gradio is a simple and easy-to-use Python library used for quickly building interactive interfaces for machine learning models. In this section, we will specifically introduce how to quickly deploy RKLLM-Server-Gradio on a Linux device using Gradio, and directly access the server within the local network to perform RKLLM model inference. The following Figure shows an example of the web interface after successfully deploying RKLLM-Server-Gradio:

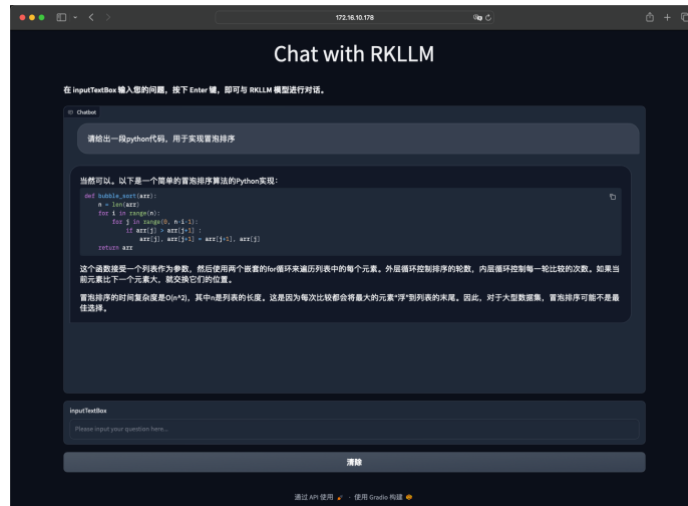


Figure 3-4 External access page for RKLLM-Server deployment example

In the current `rkllm_server_demo` directory, the specific implementation code for the RKLLM-Server-Gradio deployment example shown in Figure 3-3 is included. Users can also directly use the one-click deployment script `build_rkllm_server_gradio.sh` to quickly set up RKLLM-Server-Gradio. After successful deployment, users can choose to access the RKLLM model for inference either through the web interface or via API access.

3.4.2.1 Server-side: Introductions for RKLLM-Server-Gradio Example

The one-click deployment script `build_rkllm_server_gradio.sh` is designed to facilitate the quick setup of RKLLM-Server-Gradio on a Linux development board. Similar to the setup process for RKLLM-Server-Flask, users should pay attention to the following points before using the one-click deployment script:

1) Ensure that the development board is connected to the network via an Ethernet cable. Use the `ifconfig` command in the `adb` shell to query the specific IP address of the development board. RKLLM-Server-Gradio will be set up as a server within the local network using this IP address to accept client access.

2) Users need to complete the smooth conversion of the RKLLM model and have pushed the RKLLM model to the Linux board before executing the one-click deployment script.

The usage of the one-click deployment script `build_rkllm_server_gradio.sh` is similar to that of `build_rkllm_server_flask.sh`:

```
./build_rkllm_server_gradio.sh
--workshop [RKLLM-Server Working Path]
--model_path [Absolute Path of Converted RKLLM Model on Board] --
platform [Target Platform: rk3588/rk3576]
[--lora_model_path [Lora Model Path]]
[--prompt_cache_path [Prompt Cache File Path]]
```

Similarly, the workshop parameter specifies the working directory for RKLLM-Server-Gradio on the device; the model_path parameter indicates the absolute path of the RKLLM model on the device, which was converted using RKLLM-Toolkit, and RKLLM-Server-Gradio will read the model from this path during operation; the platform parameter specifies the platform type being used, such as rk3588 or rk3576; lora_model_path and prompt_cache_path are optional parameters, allowing the user to specify the file paths for loading a Lora model or utilizing the Prompt feature if needed.

Users can directly call the build_rkllm_server_gradio.sh script on the PC side (not on the development board) using the following command to quickly deploy the RKLLM-Server-Gradio example:

```
./build_rkllm_server_gradio.sh
--workshop /user/data
--model_path /user/data/model.rkllm
--platform rk3588
```

After executing the above command, the one-click deployment script will perform the following steps:

- 1) Check the Linux environment on the board.
- 2) Automatically install the Gradio library if not already installed.
- 3) Push the necessary files under rkllm_server_demo/rkllm_server to the board.
- 4) Index the RKLLM model in the preset working directory for RKLLM-Server-Gradio.

Once you see the message "RKLLM Model has been initialized successfully!" in the terminal, it indicates that the RKLLM-Server-Gradio example has been successfully launched.

```
warnings.warn(

=====init....=====
rkllm-runtime version: 1.0.2b9, rknpu driver version: 0.9.7, platform: RK3588
load prompt cache from '/data/cw/prompt_cache.bin'
loaded a prompt cache with prompt size of 27 tokens
RKLLM Model has been initialized successfully!
=====
Running on local URL: http://0.0.0.0:8080

To create a public link, set `share=True` in `launch()``.
```

Figure 3-5 Successful deployment of RKLLM-Server-Gradio in terminal

By referencing the specific code in build_rkllm_server_gradio.sh, users can understand the detailed deployment process of the RKLLM-Server-Gradio example. This can help users deploy custom servers

more flexibly. It is important to emphasize that in step 3 of `build_rkllm_server_gradio.sh`, the one-click deployment script automatically synchronizes the current version of RKLLM Runtime to `rkllm_server/lib/librkllmrt.so`. This ensures that `gradio_server.py` indexes `librkllmrt.so` when running, and users need to pay attention to the invocation of `librkllmrt.so` when customizing the server.

3.4.2.2 Server-side: Introductions for RKLLM-Server-Gradio Example

The deployment implementation of RKLLM-Server-Gradio is similar to RKLLM-Server-Flask. It also uses the `ctypes` library to directly call the RKLLM Runtime library to perform RKLLM model inference. The specific deployment implementation process can be referenced in Figure below:

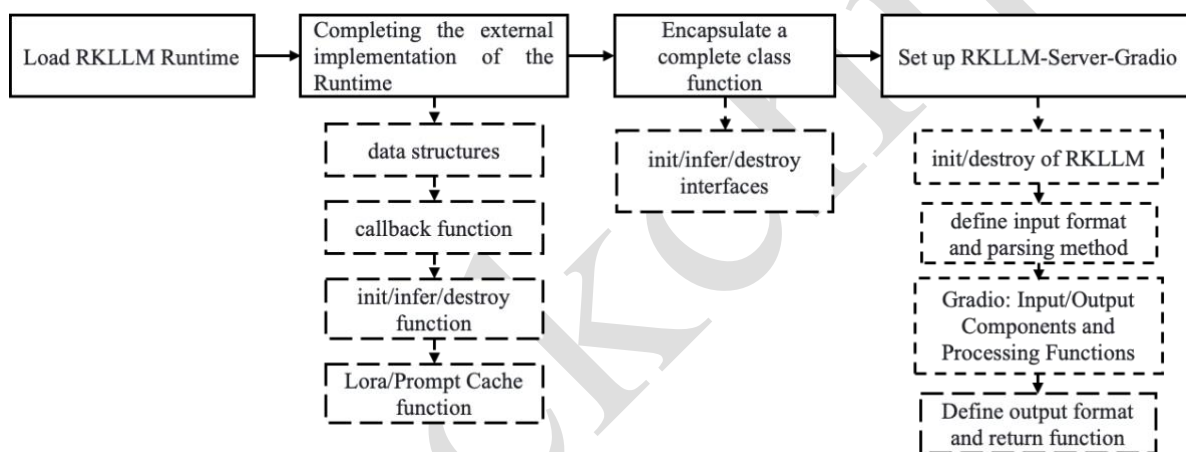


Figure 3-6 Overview of the RKLLM-Server-Flask Deployment Implementation Process

The difference is that RKLLM-Server-Gradio chooses to use the `gradio` library to implement the server setup and complete communication with the client, providing a simple web-based service. This requires specific handling of the Gradio interface in RKLLM-Server-Gradio as follows:

- 1) The Gradio function library provides complete communication for input and output data, so there is no need to define complex input-output implementations for the server via Flask.
- 2) Gradio is a deployment framework based on different control elements. During usage, it is necessary to call various Gradio components to complete the interface design and specify the trigger conditions, function call logic, and the data flow logic between different components for each element.

Users can refer to the main code in `rkllm_server/gradio_server.py` to understand the specific implementation of RKLLM-Server-Gradio, and by modifying the initialization definitions for the RKLLM model within it, they can implement different custom models. Additionally, users can refer to the deployment example of the RKLLM-Server-Gradio to deploy their own custom server.

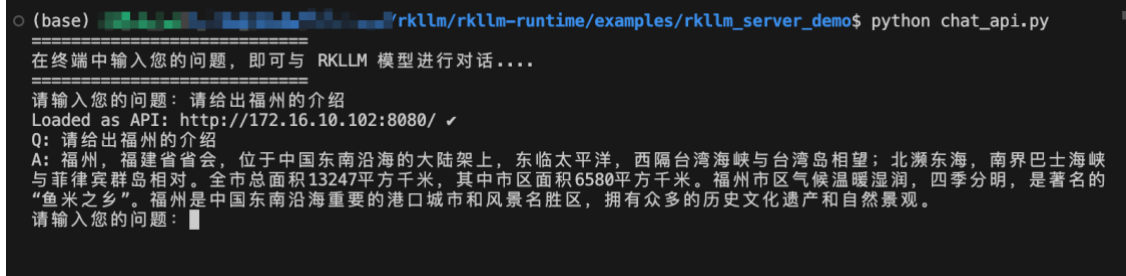
3.4.2.3 Client: RKLLM-Server-Gradio Usage Instructions

After successfully deploying the RKLLM-Server-Gradio on a Linux development board, users can access it via two methods: "Interface Access" and "API Access".

1) Interface Access: Upon successfully starting the RKLLM-Server-Gradio with the one-click deployment script, users can directly access the RKLLM model for quick interaction by opening any web browser on a computer within the current local area network and navigating to "board_IP:8080" (e.g., "172.16.10.178:8080" as shown in Figure 3-2). Gradio automatically integrates Markdown, HTML, and other syntaxes, adapting to the format of the RKLLM model's output results, such as code snippets and Markdown text. Additionally, during the setup of RKLLM-Server, an access queue is initiated. When multiple users interact with the RKLLM-Server simultaneously, the inputs are processed and returned in the order they were submitted. It's important to note that when a user's interaction with the RKLLM-Server is in the inference state (i.e., the dialogue box is highlighted), the server will not accept the user's next input until the current inference is completed.

2) API Access: In the `rkllm_server_demo` directory, `chat_api_gradio.py` is provided. After installing `gradio_client` on the PC (using the command: `pip install gradio_client`), users can interact with the RKLLM-Server solely through the API interface without relying on the graphical interface, as shown in following Figure. Before using `chat_api_gradio.py`, it's important to modify the IP address in the code to match the current IP address of the development board, as shown in the following code.

```
from gradio_client import Client  
client = Client("http://172.16.10.169:8080/")
```



```
o (base) /rkllm/rkllm-runtime/examples/rkllm_server_demo$ python chat_api.py  
=====  
在终端中输入您的问题，即可与 RKLLM 模型进行对话....  
=====  
请输入您的问题：请给出福州的介绍  
Loaded as API: http://172.16.10.102:8080/ ✓  
Q: 请给出福州的介绍  
A: 福州，福建省省会，位于中国东南沿海的大陆架上，东临太平洋，西隔台湾海峡与台湾岛相望；北濒东海，南界巴士海峡与菲律宾群岛相对。全市总面积13247平方千米，其中市区面积6580平方千米。福州市区气候温暖湿润，四季分明，是著名的“鱼米之乡”。福州是中国东南沿海重要的港口城市和风景名胜，拥有众多的历史文化遗产和自然景观。  
请输入您的问题： █
```

Figure 3-7 Access the RKLLM-Server-Gradio via API calls in terminal

Users can choose between the two client invocation methods based on their specific needs. For instance, when providing interactive services within a local area network, it's recommended to use the interface access method. On the other hand, if customizing access behaviors to RKLLM-Server-Gradio is required, it's advisable to use API Access for further development.

Lastly, it's important to note that in the implementation of RKLLM-Server-Gradio, there isn't a definition of data structures for sending and receiving data similar to OpenAI-API. Therefore, this deployment implementation is not compatible with the OpenAI-API interface. When conducting further development, users should refer to the specific function implementation in `chat_api_gradio.py`. If compatibility with the OpenAI-API interface is required, please refer to the implementation of RKLLM-Server-Flask.

4 Reference

RKLLM: <https://github.com/airockchip/rknn-llm>

RKNN: <https://github.com/airockchip/rknn-toolkit2>

Rockchip